Marcus Vance

marcus.vance@example.com | (555) 432-8765 | San Francisco, CA | github.com/marcusvance | linkedin.com/in/marcusvance

PROFESSIONAL SUMMARY

Senior Full-Stack AI Software Engineer with 6 years of production experience architecting scalable distributed systems, web applications, and agentic workflows. Deep expertise in React 18, Node.js (ES modules), Express, MongoDB Atlas, and runtime schema enforcement with Zod. Proven track record deploying LangGraph state machines, containerizing microservices via Docker, and achieving 90%+ test coverage with Vitest.

PROFESSIONAL EXPERIENCE

NexusCloud AI - Staff Software Engineer (2022 - Present)

- Architected enterprise multi-agent workflow engine using LangGraph, Node.js, and Express handling 4M monthly requests.
- Enforced strict runtime type boundaries across all LLM pipeline inputs and outputs using Zod schema validation.
- Optimized high-throughput MongoDB Atlas database clusters, reducing aggregate query execution time by 44%.
- Built modular frontend dashboards with React 18, Vite, and modern CSS modules with real-time WebSocket state streaming.
- Established CI/CD pipelines utilizing Docker multi-stage builds and automated Vitest test suites with 92% coverage.

Apex Digital Systems - Senior Full-Stack Developer (2018 - 2022)

- Led development of core MERN stack SaaS platform from inception to acquisition with $12M ARR.
- Built secure REST APIs in Express with rate limiting, Helmet, CORS, and JWT authentication.
- Implemented RAG search pipelines with vector embeddings and cosine similarity index tuning.

EDUCATION

University of Washington - B.S. in Computer Science (2014 - 2018) | Magna Cum Laude, GPA 3.91/4.0

TECHNICAL SKILLS

Languages: TypeScript, JavaScript (ES6+), Python, SQL | Frontend: React, Vite, Next.js, CSS3 | Backend: Node.js, Express, REST APIs, WebSockets | AI/ML: LLM APIs (Gemini, OpenAI), LangGraph, RAG, Embeddings | Databases & Tools: MongoDB, PostgreSQL, Zod, Docker, Vitest, Jest, Git